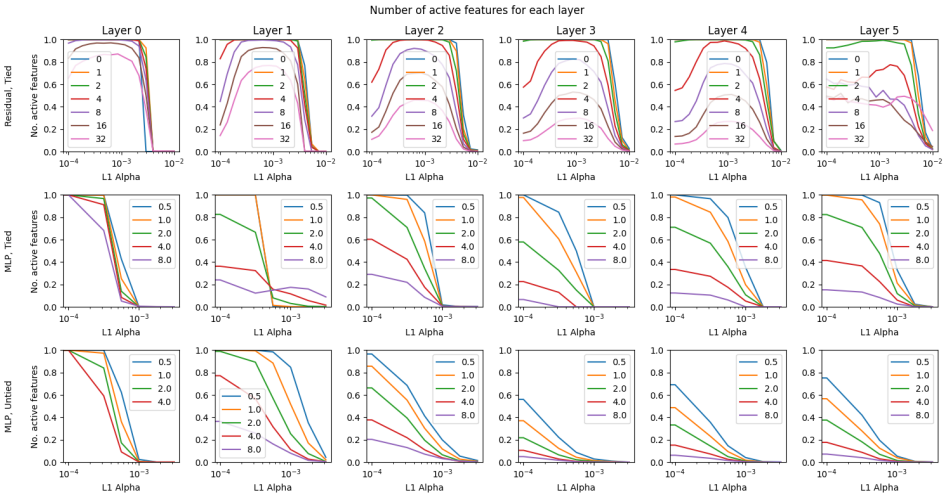

Figure 16: The number of features that are active, defined as activating more than 10 times across 10M datapoints, changes with sparsity hyperparamter α and dictionary size ratio R.

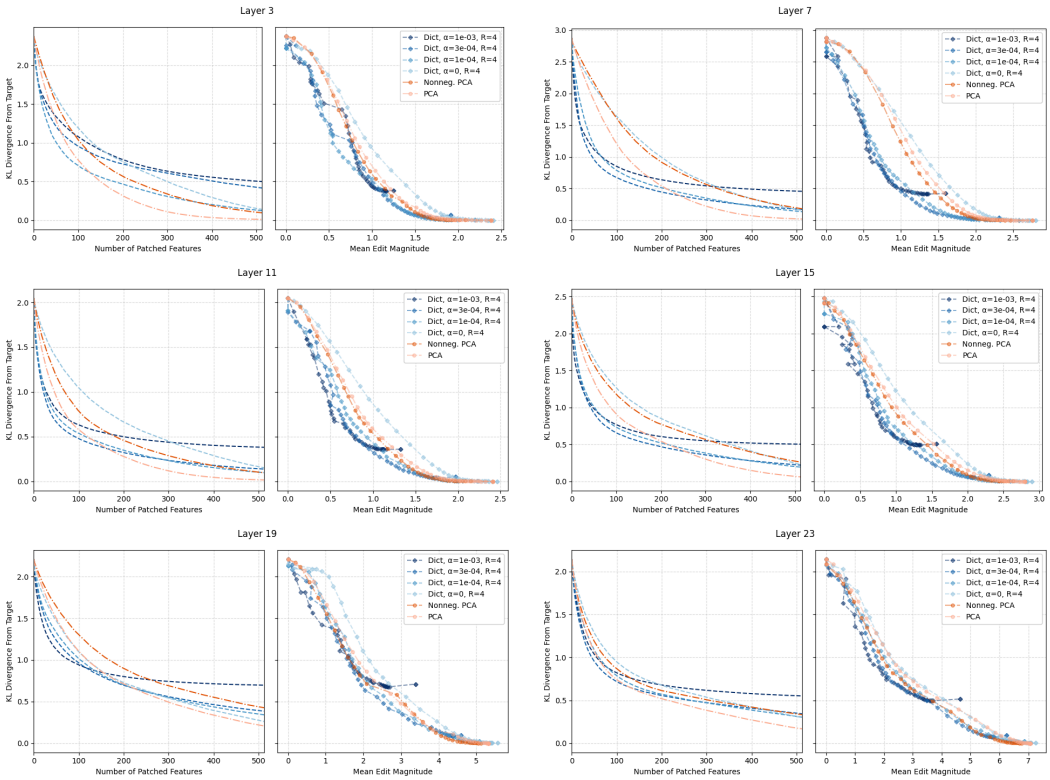

Figure 17: Divergence from target output against number of features patched and magnitude of edits for layers 3, 7, 11, 15, 19 and 23 of the residual stream of Pythia-410M. Pythia-410M has 24 layers, which we index 0, 1, ..., 23.

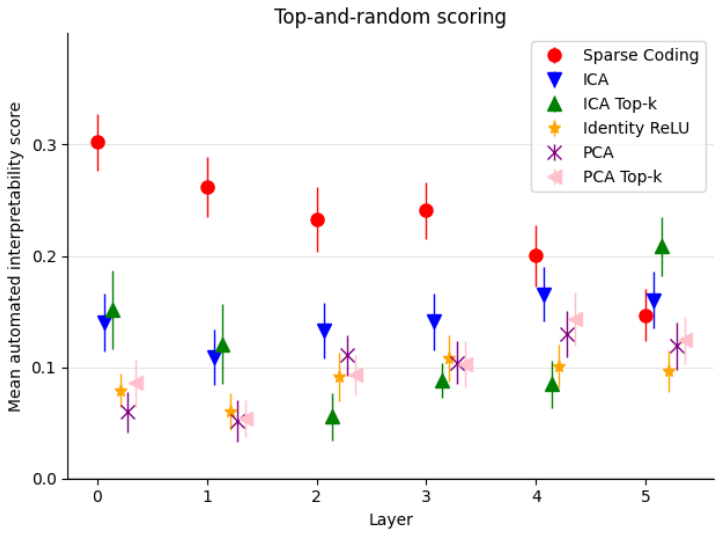

Figure 18: Autointerpretation scores across layers for the residual stream, including top-K baselines for ICA and PCA.